

AI CHAMPIONS · FUNDAMENTALS

Prompt Engineering Fundamentals

How to give AI clear instructions, relevant information, and usable output requirements — in one session.



WHY THIS MATTERS

Prompt quality decides output quality

An unclear request gets a generic, unreliable response — something you end up rewriting by hand. A clear request gets a relevant, usable, reviewable response — something you can check and send.

The reason is simple, and it is the core idea of this session: the model cannot infer important business context that you have not provided. It only knows what you tell it.

THE PATTERN

X Unclear ask

→ generic, unreliable, reworked by hand.

✓ Clear ask

→ relevant, usable, reviewable.

This session is about moving from the first line to the second — on your real work.



LEARNING OBJECTIVES

By the end of this session, you will be able to:

1

Explain what a prompt is

The parts that guide an AI response — and which of them you control.

2

Build a clear prompt with CO-STAR

Context, Objective, Style, Tone, Audience, Response.

3

Provide source material and examples

Separate instructions from material; ground it; show one done right.

4

Improve a response with follow-ups

Refine and diagnose instead of restarting.

5


Review before you use it

Check accuracy and usefulness — and protect member data.



THE JOURNEY

Four stages — and what today covers



Keep an eye on the top-right corner of every slide — it shows where you are on this journey.



THE BASICS

What is a prompt?

A prompt is the combination of instructions, context, source information, examples, and output requirements you provide to guide an AI response.

It is not just the sentence you type. Everything you give the model shapes what comes back — and everything you leave out, it has to guess.

A PROMPT CAN INCLUDE

- Typed instructions
- Pasted text
- Uploaded documents
- Examples of what you want
- Earlier turns in the same conversation
- Approved connected information



STAGE 1 · ASKING

Where everyone starts

The most common prompt in the building — not wrong, just incomplete. Today we build the prompt that fixes it, on one real batch of feedback.

▶ THE ONE-LINE PROMPT

Summarize this feedback.

WHAT COMES BACK

“Members had mixed feedback about their visit — some positive comments and some complaints. Overall it was a fairly typical weekend.”

No audience. No purpose. No format. No priorities. No guard against invented claims.

OUR RUNNING EXAMPLE

What it was asked to summarize — last weekend's survey, about 30 comments. Five of them:

“Waited 22 min at checkout — only 3 lanes open.”

“Rotisserie chicken is always fresh!”

“The app froze when I clipped the gas coupon.”

“Gas pump 4 is down again.”

“Wilted spinach and moldy berries in produce.”

A safety issue, a staffing problem, a broken app — all flattened into “mixed feedback.”



STAGE 2 · COMPOSING

CO-STAR: a practical checklist for a strong prompt

**C — Context**

Relevant background and source information.

**O — Objective**

What you want it to accomplish.

**S — Style**

How it should be written or presented.

**T — Tone**

The voice it should use.

**A — Audience**

Who will read or use the output.

**R — Response**

The exact output structure you need.

Not every prompt needs every element — use the ones that reduce ambiguity.



STAGE 2 · COMPOSING

Building the prompt — element by element



Each element joins the prompt →



Context

The background only you have — the situation, the material, the time period. It cannot know this unless you say it.

✗ “Analyze this.”

THE PROMPT SO FAR

You are reviewing 30 anonymized comments from last weekend's club survey.



STAGE 2 · COMPOSING

Building the prompt — element by element

C**O**

S

T

A

R

Each element joins the prompt →

Objective

State the task as an action: summarize, compare, classify, prioritize, extract. Unnamed, the model picks one for you.

✗ *“Help me with this feedback.”*

THE PROMPT SO FAR

You are reviewing 30 anonymized comments from last weekend's club survey. **Group the comments into themes and count each.**



STAGE 2 · COMPOSING

Building the prompt — element by element



Each element joins the prompt →



Audience

Name who will use the answer — it quietly decides style, tone, depth and length. The highest-leverage single sentence.

✗ (no reader named)

THE PROMPT SO FAR

You are reviewing 30 anonymized comments from last weekend's club survey. Group the comments into themes and count each. **Write for a busy club manager.**



STAGE 2 · COMPOSING

Building the prompt — element by element



Each element joins the prompt →

S·T Style & Tone

The shape (bullets, summary, table) and the voice (neutral, urgent). The audience you just named implies both — say it anyway.

✗ *“Make it sound professional.”*

THE PROMPT SO FAR

You are reviewing 30 anonymized comments from last weekend's club survey. Group the comments into themes and count each. Write for a busy club manager. **Keep it concise and neutral — no technical terms.**



STAGE 2 · COMPOSING

Building the prompt — element by element

C

O

S

T

A

R

Each element joins the prompt →

R

Response

The exact output: how many bullets, what each must include, what comes first. Everything here is checkable later.

✗ *“Summarize the feedback.”*

THE PROMPT SO FAR

You are reviewing 30 anonymized comments from last weekend's club survey. Group the comments into themes and count each. Write for a busy club manager. Keep it concise and neutral — no technical terms. **Return at most five bullets, and put any potential safety issue first.**

**STAGE 2 · COMPOSING — THE PAYOFF**

Run it — the prompt you just built

One canvas, six elements, one prompt — assembled click by click in front of you, plus one guard line at the end that the next stage unpacks. Same batch, same model; the only thing that changed is what you told it.

WHERE WE STARTED (SLIDE 6)**THE PROMPT**

Summarize this feedback.

WHAT YOU GET BACK

Members had mixed feedback about their visit — some positive comments and some complaints. Overall it was a fairly typical weekend.

WHY IT FALLS SHORT

You gave it six words — no role, no reader, no shape to aim at. It hedged because guessing was all it could do.

THE PROMPT YOU BUILT**THE PROMPT**

You are reviewing 30 anonymized comments from last weekend's club survey. Group the comments into themes and count each. Write for a busy club manager. Keep it concise and neutral — no technical terms. Return at most five bullets, and put any potential safety issue first. Do not infer facts not stated in the comments.

WHAT YOU GET BACK

- Safety: gas pump 4 down — 1.
- Checkout waits — 12.
- Produce quality — 7.
- Fresh food — 5 (+).
- App issues — 4.

WHY IT WORKS

Each element handed over a fact only you knew — and the guard line closed the door on invented claims.



CHECKPOINT · QUICK CHECK

Which elements are missing?

Two minutes, with a neighbour: check this prompt against CO-STAR before we move on.

▶ THE PROMPT TO CHECK

“Write a short summary of the produce complaints for me.”

Go letter by letter: C? O? S? T? A? R?

CHECK IT AGAINST CO-STAR

- C** — which comments? from when? how many?
- O** — “summary” is an action... of what, for what decision?
- S / T** — shape and voice unspecified.
- A** — “for me”: who are you, and who reads it next?
- R** — “short”: how short? bullets? what must it include?

The takeaway: a prompt can look reasonable and still leave almost everything to guesswork. The checklist finds the gaps in seconds — that's why it's a checklist.



STAGE 3 · DIRECTING — FIRST STEPS

Your prompt is built — now reality pushes back

**Messy input**

Long walls of text blur your task into the material → label the parts.

**Missing facts**

It fills gaps with confident guesses → ground it in the source.

**Format drift**

Descriptions of a format get ignored → show an example.

**Judgment calls**

It can't guess your priorities → give decision criteria.

Four situations, four moves — one slide each. This is the start of Directing; the Intermediate module goes deeper on every one.



STAGE 3 · DIRECTING — SEPARATE TASK FROM MATERIAL

Label what's instruction and what's material

When the input is long or messy, your task can blur into the material — so label the parts. Plain headings work: TASK, REQUIREMENTS, COMMENTS, OUTPUT. The exact words matter less than the separation.

ONE BIG BLOB

THE PROMPT

Group these into themes and count them. waited 22 min at checkout... chicken always fresh... pump 4 down... (+25 more)

WHAT YOU GET BACK

Sure — you'd like me to group and count these. The first few mention checkout waits and the rotisserie chicken, which sound positive overall...

WHY IT FALLS SHORT

With no boundary, your task and the comments blur into one wall of text — it can't tell orders from material.

CLEARLY LABELLED

THE PROMPT

*TASK — group the comments into themes; count each.
REQUIREMENTS — flag safety concerns; do not infer what isn't stated.
COMMENTS — (the 30 pasted here)
OUTPUT — theme, count, example quote, suggested follow-up.*

WHAT YOU GET BACK

Gas pump — 1 — SAFETY — escalate today
Checkout waits — 12 — 'only 3 lanes open' — staffing review
Produce quality — 7 — 'wilted spinach' — check rotation
App issues — 4 — 'froze on the coupon' — report to IT
Fresh food — 5 (+) — 'always fresh' — share the win

WHY IT WORKS

The headings draw the line between orders and material — it parses cleanly; your review still owns correctness.



STAGE 3 · DIRECTING — GROUND IT

Only the source, quoted — or “not stated”

When accuracy beats fluency — vendor terms, policies, anything with numbers — hand it the real document and close every exit: answer only from this, quote the line, and say “not stated” rather than guess.

FROM MEMORY**THE PROMPT**

What are this vendor's delivery terms?

WHAT YOU GET BACK

Typically 3–5 business days, with free returns within 30 days.

(None of this appears in the document.)

WHY IT FALLS SHORT

Asked from memory, it fills the gaps with plausible patterns — fluent, confident, and never checked against anything.

GROUNDED IN THE DOC**THE PROMPT**

Using only the vendor one-pager below, list the delivery terms. Quote the relevant line for each. If a term is missing, write 'not stated'.

WHAT YOU GET BACK

- 'Orders ship within 10 business days.'
- 'Freight billed at cost.'
- Return window — not stated in the document.

WHY IT WORKS

'Only the document, quote it, or say not stated' leaves nowhere to invent — and makes your review a two-minute check, not a hunt.



STAGE 3 · DIRECTING — SHOW AN EXAMPLE

When describing the format isn't enough

When a format is hard to describe — or consistency across many items matters — show one or two done right and say “match these.” It copies the pattern faithfully. It copies mistakes just as faithfully, so check your examples first.

TOLD IT — “format these neatly”

THE PROMPT

Format these comments neatly.

WHAT YOU GET BACK

- checkout was really slow saturday, 3 lanes open

The rotisserie chicken is great, always fresh!

- app froze when I clipped the gas coupon

Gas pump 4 down again, ugh.

wilted spinach and moldy berries in produce

WHY IT FALLS SHORT

'Neatly' is a wish, not a format — with nothing to copy, it picked a different shape for every row.

SHOWED IT — two examples first

THE PROMPT

Here are two comments classified correctly:
Checkout waits — Negative — 'slow.'
Rotisserie chicken — Positive — 'fresh.'
Apply the same structure to the remaining comments.

WHAT YOU GET BACK

Checkout waits — Negative — 'only 3 lanes open.'

Rotisserie chicken — Positive — 'always fresh.'

App issues — Negative — 'froze on the gas coupon.'

Gas pump — Negative — 'pump 4 down again.'

Produce — Negative — 'wilted spinach, moldy berries.'

WHY IT WORKS

Two finished rows define the format better than a paragraph of rules — it simply continued the pattern it could see.



STAGE 3 · DIRECTING — GIVE DECISION CRITERIA

Tell it what matters before it judges

When a task needs judgment — ranking, escalating, choosing — name the factors that count and ask for an evidence-based reason per call. Your priorities cannot be guessed; they have to be stated.

NO CRITERIA

THE PROMPT

Which three issues should we escalate?

WHAT YOU GET BACK

1. Checkout waits (most comments)
2. Produce quality
3. App issues

(Picks the most-repeated. The gas pump — a safety issue mentioned once — doesn't make the list.)

WHY IT FALLS SHORT

With no factors named, it fell back on the loudest signal — repetition — which is not the same as importance.

CRITERIA STATED

THE PROMPT

Rank the issues by: 1) safety risk, 2) members affected, 3) urgency, 4) effort to resolve. Give a one-sentence, evidence-based reason for each ranking.

WHAT YOU GET BACK

1. Gas pump 4 — safety risk; 'down again' suggests a repeat failure.
2. Checkout waits — 12 members affected; staffing is actionable.
3. App issues — 4 members; a ticket for IT.

WHY IT WORKS

Your priorities entered the decision, and the evidence-based reasons let you check every ranking in seconds.



IMPROVE · REFINE, DON'T RESTART

Improve it in one follow-up

Even a well-built prompt rarely lands perfectly first time. Don't start over — everything you built stays in the conversation. Name what's off, change exactly that, and keep the rest.

FIRST TRY

THE PROMPT

Summarise this week's feedback in a short note to the manager.

WHAT YOU GET BACK

This week saw complaints about checkout waits and produce, plus praise for the fresh food. A few app and gas-pump issues too.

WHY IT FALLS SHORT

A first prompt rarely names everything you care about — you only see what's missing once you read an answer.

ONE REFINEMENT

THE PROMPT

Good start. Convert it into five bullets and put the safety issue first.

WHAT YOU GET BACK

- Safety: gas pump 4 down.
- Checkout waits — top complaint.
- Produce quality.
- App issues.
- Wins: fresh food.

WHY IT WORKS

The follow-up kept all the context you'd built and changed only what you named — far cheaper than starting over.



IMPROVE · DIAGNOSE THE MISS

Name the problem, then fix that one thing

Too generic

→ Add the purpose, the audience, and the detail you need.

Wrong format

→ Specify the structure and the length.

Unsupported claims

→ Restrict it to the provided sources.

Missing priority

→ Add your decision criteria.

Too verbose

→ Set a limit and name the audience.

Important item omitted

→ State what must be included.

Inconsistent output

→ Provide an example to copy.



REVIEW · VERIFY BEFORE USING

A strong prompt doesn't replace your review

**Facts supported?**

Every claim traces to the source material.

**Counts correct?**

Numbers and totals actually add up.

**Citations real?**

Quoted lines exist in the source.

**Instructions followed?**

Format, length, and priorities as asked.

**Anything missing?**

The important item didn't get dropped.

**Right for the reader?**

Language fits the audience — and nothing restricted.



REVIEW · RESPONSIBLE PROMPTING

Approved tools, approved information

Use only approved AI tools and approved information. Do not enter restricted, confidential, personal, payment, or sensitive member information unless the tool and the use case are explicitly approved for that data.

And remember what a prompt is not: a prompt instruction is not a substitute for company policy, access controls, security review, or your own judgment.

THREE SAFEGUARDS

- 1** Use approved, properly handled source material.
- 2** Tell the model not to infer unsupported facts.
- 3** Review the result before you share or act on it.



PRACTICE · YOUR TURN

Rebuild this prompt with CO-STAR

Everything from today, on one real task. Work in pairs against the practice dataset — then compare prompts before you compare outputs.

▶ THE WEAK PROMPT

“Summarize these comments for leadership.”

Rebuild it: add the context, the objective, the reader, the style and tone, and the exact output — then add the guard line.

Material: the 29-comment practice dataset handout.

A STRONG ANSWER INCLUDES

- C** — the batch, named: which survey, roughly how many.
- O** — an action: group into themes and count.
- A** — the leader who will actually read it.
- S/T** — concise, neutral, plain language.
- R** — bullets with counts, safety first.
- +** — the guard line: nothing invented, nothing restricted.



RECAP · TAKE IT WITH YOU

Everything we promised, on one card

THE JOB AID

C Context · **O** Objective · **S** Style · **T** Tone · **A** Audience · **R** Response

After prompting — three habits:

Ground it. Refine it. Verify it.

YOUR JOURNEY

Today you covered **Asking** → **Composing**, and took your first steps into **Directing**.

Next stops when you're ready:

Intermediate — Directing, in depth.

Advanced — Designing: chains, templates, evaluation.

This week: use CO-STAR on one real, approved task — then ground it, refine it, and verify it before it goes anywhere.