



BJ'S AI ENABLEMENT · CORE COURSE

Thinking with AI

How to decide where and how AI fits into your work — before the prompting starts. For ChatGPT and Claude, on real BJ's work.



LEARNING OBJECTIVES

By the end of this module, you will be able to:

1

Decide if and how AI fits a task

Three quick tests, then five questions that make the task clear.

2

Choose the role AI should play

Seven roles — and why that choice matters more than the wording.

3

Check AI's output before you act on it

Five checks to run before anything ships.

4

Run the full loop on your own work

Think → Partner → Decide, in about two minutes.



WHY THIS COURSE

AI value doesn't come from typing better prompts

Having access to ChatGPT and Claude is not the same as knowing how to use them on real BJ's work.

The value comes from rethinking the work itself: getting clear on the outcome, deciding what AI should help with, and applying your own judgment before anything ships.

Research backs this up: in a large study of AI users at work, the more people trusted the AI, the less critical thinking they applied to its output. This course is the countermeasure. Prompt wording matters too — and it has its own course next.

OUR RUNNING EXAMPLE

Maria manages a BJ's club. Her regional manager asked her to propose how her club can improve membership renewals.

One real task — we'll follow it through the whole course.



THE PROBLEM — AND WHERE WE'RE GOING

Two people, same tool, same task — different results

TYPE FIRST, THINK NEVER

"Give me ideas to improve membership renewals."

WHAT YOU GET

Five generic bullets any retailer could get: better service, reminder emails, loyalty rewards...

Nothing about her club, her members, or what's already been tried.

THE SAME PERSON, AFTER THIS COURSE

Same task — but she clarified the outcome, chose AI's role, and checked the output.

WHAT YOU GET


Three specific actions for her club's next 90 days — with the weak spots already found.

How? That's the next 25 minutes — three steps, nothing magic.



THE HABIT

The AI Thinking Loop



Prompting starts after you understand the task, the role AI plays, and what stays yours.



STEP 1 · THINK — FIRST, WHETHER

Should AI touch this task at all?



Is it high-stakes and hard to undo?

A termination, a member dispute, a safety call — some decisions are yours alone.



Is your judgment the whole point?

If your reasoning and your name are the deliverable, AI can inform it — not make it.



Does the data have rules?

Member, payroll, or confidential data — know BJ's data guidance before you paste.

Maria's proposal clears all three in ten seconds. When a test says yes, the task stays yours — for everything else, run the loop.



STEP 1 · THINK

Get clear on the task — before you open AI



What outcome am I trying to produce?

A decision? A proposal? Three options? Name it.



Who is this for?

Your regional manager reads differently than your team.



What context matters?

What only you know: your club, your members, what's been tried.



What does good look like?

Specific, doable, defensible — define it before you start.



Which parts need my judgment?

What must come from you, no matter what AI says.



STEP 1 · IN ACTION

Thirty seconds of thinking changed the task

Maria ran her assignment through the five questions before opening ChatGPT or Claude.

“Improve membership renewals” became: “Propose three actions my club can take in the next 90 days to lift renewals — for my regional manager, drawing on what I know about my members and what we tried last quarter.”

Vague task in, vague answers out. Clear task in, and everything downstream gets better — no clever wording required.

THE TASK, CLARIFIED

Outcome: 3 actions, 90 days

For: her regional manager

Context: renewal dip since spring;
reminder-call push already tried

Her judgment: which actions her team can
actually run



STEP 2 · PARTNER

Decide what role AI should play

HELP ME SEE

Organize information

Turn scattered notes into structure.

Analyze patterns

What's changed? What stands out?

HELP ME SHAPE

Clarify the problem

"Ask me questions before you answer."

Generate options

Eight ideas beat one first draft.

Improve a draft

Tighten what you wrote — keep it yours.

HELP ME TEST

Challenge assumptions

"Where does this plan fall apart?"

Simulate a stakeholder

"Read this as my regional manager."

Most people jump straight to "write it for me." Often the task needs a different role entirely — and choosing the role is the step that matters most.



STEP 2 · IN ACTION

Same task — the role choice made the difference

WRONG ROLE: “WRITE IT FOR ME”

“Write a proposal to improve my club's membership renewals.”

WHAT YOU GET

A polished-sounding proposal built on guesses — it doesn't know her club, so it invents.

Sounds finished. Isn't hers.

RIGHT ROLE: GENERATE OPTIONS

“Here's my situation — the renewal dip, what we've tried, my weekend staffing. Give me eight options for the next 90 days.”

WHAT YOU GET

Options grounded in her club's reality — she picked three worth taking further.

AI helped her think. She did the choosing. (The pressure test comes next.)



STEP 2 · THE ROLES NOBODY USES

AI agrees with you by default — make it push back

The two most valuable roles are the two people use least: challenge assumptions and simulate a stakeholder.

ChatGPT and Claude are built to be agreeable. Left alone, they will tell you your plan is great. Real opposition only shows up when you ask for it — and that is exactly when a plan gets stronger.

One sentence turns it on. Use it before your work goes to anyone whose opinion matters.

MARIA'S MOVE

*“Act as my skeptical regional manager.
Where does this plan fall apart?”*

It flagged that her strongest action depends on weekend front-desk staffing she doesn't have.

Better to find out now than in the meeting.



STEP 2 · IN ACTION — THE DRAFT

Now — and only now — she asks it to write

The pushback did its job: Maria swapped the staffing-dependent option for another from her eight.

Then — and only then — the role most people start with: “Turn my three options into a one-page proposal draft for my regional manager.”

“Write it for me” was the wrong opening move earlier. After the thinking, it's the right closing one — this draft is built from her options, her context, her constraints.

It comes back reading clean, confident, and finished. Hold that thought.

THE CONVERSATION SO FAR

1. Context + “give me eight options”
2. Picked three worth taking further
3. “Where does this fall apart?” → swapped one
4. “Turn them into a one-page draft”



STEP 3 · DECIDE

Judgment before action — check the output



Is it accurate?

Verify names, numbers, claims. AI states guesses confidently.



Is it complete?

What did it leave out that you would be accountable for?



Is it right for BJ's?

Our members, our tone, our policies — not generic retail.



What did AI assume?

Ask it directly: “What assumptions did you make?”



What do I own?

The call, the numbers, the name on it — those are yours.



STEP 3 · THE CONFIDENCE TRAP

The more you trust AI, the less you check it

Researchers at Microsoft and Carnegie Mellon surveyed 319 professionals who use AI at work. The finding: the more confident people were in the AI, the less critical thinking they applied to its output. Confidence in their own expertise had the opposite effect — more scrutiny, better results.

So the risk isn't beginners. It's experienced users on their hundredth task, when checking starts to feel unnecessary.

That's why the five checks are a rule, not a mood. Run them especially on the days you're sure you don't need to.

THE TRAP IN ONE LINE

As trust in the tool goes up, scrutiny goes down.

The checks matter most exactly when they feel least necessary.

Maria's draft is about to be exactly this moment.



STEP 3 · IN ACTION

It looked shippable — the checks said otherwise

Maria's draft read clean and confident. Two of the five checks caught problems anyway.

Is it accurate? One option offered a renewal discount — not a club-level call. Plausible, confident, wrong.

What did AI assume? Asked directly: that most members auto-renew. At her club, most don't — which changes which actions matter.

The loop-back: “Drop the discount — clubs can't set membership pricing. Assume most members do not auto-renew. Rework it.” One more pass, one more check — clean.

THE STANDARD

The output isn't done when AI stops typing. It's done when you can defend it.

If you can't defend it, don't ship it.

And a failed check isn't the end — take it back to the conversation with what you now know.



THE FULL LOOP

The whole habit on one real task — about two minutes

1

1. THINK — 30 seconds

Outcome: 3 actions, 90 days. For: her regional manager. Context: renewal dip; reminder calls already tried. Hers: the final call.

2

2. PARTNER — 1 minute

Generate options (“give me eight”) → pressure-test with her skeptical regional manager → swap one option → turn the three into a one-page draft.

3

3. DECIDE — 30 seconds, plus one loop back

Caught the discount error and the auto-renew assumption — took both back to the conversation, reworked, re-checked. Kept three actions she can defend.

Result: a proposal that is specific to her club, tested against pushback, and genuinely hers.



WHAT YOU CAN NOW DO

You have met every goal we set at the start



You decide if AI fits, then clarify

Three quick tests, five questions — before AI is involved at all.



You choose AI's role

See, shape, or test — not just “write it for me.”



You check before you act

Five checks. If you can't defend it, don't ship it.



You can run the loop

Think → Partner → Decide, on any task, in about two minutes.



YOUR TURN

One task this week

Before you leave: write down one real task from this week — then which gate test or which of the five questions it fails today. That's your first rep. Then join us for Prompt Engineering Essentials: now that you know where AI fits, learn how to get the most from every exchange.

OPTIONAL — GO DEEPER

Appendix

One optional page — where this course comes from, for anyone who wants to check our work.



OPTIONAL

Where this course comes from

The structure of this course follows published work; the specific questions, roles, and checks were built at BJ's on that foundation.

The three steps align with the AI Fluency framework published by Anthropic with University College Cork and Ringling College: Delegation (our Think), Discernment and Diligence (our Decide). Their fourth competency — Description, the wording of prompts — is our next course.

The confidence trap and the three-phase picture of critical thinking with AI come from a 2025 Microsoft Research and Carnegie Mellon survey of 319 knowledge workers.

GO DEEPER

AI Fluency: Framework & Foundations —
free, about 3 hours
anthropic.com/ai-fluency

“The Impact of Generative AI on Critical
Thinking” (Lee et al., 2025)
microsoft.com/research